



An End-to-End Workflow for Data-Driven GPU Optimization with LLVM


LLVM @ CGO26
Jan. 31st 2026

Konstantinos Parasyris

Center for Advanced Scientific Computing (CASC), LLNL


Prepared by LLNL under Contract DE-AC52-07NA27344.

How Science Actually Advances





Progress comes from *fast, controlled iteration* and *trustworthy observation*.

How Science Actually Advances



Reducing the complexity and time of this cycle is a strong indicator of continuous scientific progress

How Science Actually Advances



LLVM's Strength — and Its Accessibility Problem


LLVM's Strength



Rich observability & optimization power

Rich observability & optimization power

LLVM's Accessibility Problem



Steep learning curve

Compiler-centric tooling

Manual kernel extraction

Fragile scripts

Long iteration cycles


High barrier for non-compiler experts




Mneme is a tool providing a core infrastructure to enable data driven research on GPU LLVM research

Mneme

What is Record Replay?




What is Record Replay?







What is Record Replay?



Record Replay can provide



Mneme: Record–Replay for Scalable Optimization and analysis



- Implements record–replay for GPU
- Decouples tuning from application dependencies
- Integrates with LLVM
 - Python accessors to Functions, Blocks, Instructions etc.
 - Similar to numba/llvmlite
 - Proteus is the execution engine and applies optimizations
- Exposes replayed kernels to Python ecosystem
- Enables autotuning, analysis, and experimentation

High Level Of Execution Phases



Build “Mneme”

- `export LLVM_INSTALL_PATH=${ROCM_PATH}`
- `pip install https://github.com/olympus-HPC/Mneme`



Create a “recordable executable”

- Apply instrumentation pass to the code



Record the execution of an application

- Check the generated artifacts



Replay a single Kernel

- Verify outputs
- Create your own autotuner

Create a “recordable executable”

1

Include Mneme on build process

```
> cat CMakeLists.txt  
...  
find_package(HIP REQUIRED)  
find_package(mneme REQUIRED)  
add_executable(tutorial.exe tutorial.hip)  
add_mneme(tutorial.exe)  
...
```

2

Configure & Build

```
> cmake -B BUILD -S SRC_PRJ \  
-DCMAKE_C_COMPILER=$(mneme config cc) \  
-DCMAKE_CXX_COMPILER=$(mneme config cxx) \  
-DCMAKE_PREFIX_PATH=$(mneme config cmakedir)  
> cmake --build BUILD/
```



The executable carries its own compiler IR



What:

- Embedded LLVM IR

Why it matters:

- Enables **post-mortem analysis and recompilation**
- No need to recover IR from build system or source tree



All kernel executions become observable and interceptable

What:

- Kernel launches go through **Proteus API**
- Vendor launch APIs are not invoked directly

Why it matters:

- Intercept kernel launches, arguments, launch configurations etc.
 - These can be “tunable parameters” at replay time


Record the execution of an application

1

Wrap “recordable executable” execution with mneme

```
> mneme record -rdb record-db-dir/ -vass X \
   -- <recordable-executable> \
   <arguments>
```

Trace of host-device events



- 1) Store Mneme Memory to persistent storage (prologue)
- 2) Query proteus for LLVM IR of the kernel and store into storage
- 3) Launch Kernel (synchronously)
- 4) Store Mneme Memory to persistent storage (epilogue)

Address Space
Managed by Mneme
(|HighAddress – LowAddress| = “vass”)



Record the execution of an application

1

Wrap “recordable executable” execution with mneme

```
> mneme record -rdb record-db-dir/ -vass x \
   -- <recordable-executable> \
   <arguments>
```

2

Recording artifacts are stored under “record-db-dir”

```
> tree record-db-dir/
  └── <static-hash>.json
  ├── DeviceState.epilogue.<static-hash>.<dynamic-hash>.mneme
  ├── DeviceState.prologue.<static-hash>.<dynamic-hash>.mneme
  └── RecordedIR_<static-hash>.bc
```

Replay a single Kernel

1) Replay a single kernel invocation

```
> mneme replay \
    -rdb record-example-dir/<static-hash>.json \
    -rid <dynamic-hash> "default<03>"
```

Trace of host-device events

- ```
graph TD; A(()) --> B(()); B --> C(()); C --> D(());
```

Instantiate Device Memory Space

Initialize Memory

Compile and execute code through Proteus


Compare device memory with epilogue



(|HighAdress – LowAdress| = “vass”)



## Automated verification



# Replay a single Kernel

1

Replay a single kernel invocation

```
> mneme replay \
 -rdb record-example-dir/<static-hash>.json \
 -rid <dynamic-hash> "default<03>"
```

2

Execution emits a key-value dictionary describing various metrics

```
"Replay-config": {
 "grid": {
 "x": 40000,
 "y": 1,
 "z": 1
 },
 "block": {
 "x": 256,
 "y": 1,
 "z": 1
 },
 "shared_mem": 0,
 "specialize": false,
 "set_launch_bounds": false,
 "max_threads": null,
 "min_blocks_per_sm": 0,
 "specialize_dims": false,
 "passes": "default<03>",
 "codegen_opt": 3,
 "codegen_method": "serial",
 "prune": true,
 "internalize": true
},
```

```
"Result": {
 "preprocess_ir_time": 9.2298723757267e-06,
 "opt_time": 0.006206092890352011,
 "codegen_time": 0.01226967596448958,
 "obj_size": 4792,
 "exec_time": [
 84040,
 82561,
 81761,
 83360,
 76520
],
 "verified": true,
 "executed": true,
 "failed": false,
 "start_time": "",
 "end_time": "",
 "gpu_id": 0,
 "const_mem_usage": -1,
 "local_mem_usage": 0,
 "reg_usage": 12,
 "error": ""
}
```

# Replay a single Kernel

1

## Replay a single kernel invocation

```
> mneme replay \
 -rdb record-example-dir/<static-hash>.json \
 -rid <dynamic-hash> "default<03>"
```



*These parameters can be modified*

By forming valid configuration ranges of these parameters one can search the space and tune the application in respect to some quantity of interest

2

Execution emits a key-value dictionary describing various metrics

```
"Replay-config": {
 "grid": {
 "x": 40000,
 "y": 1,
 "z": 1
 },
 "block": {
 "x": 256,
 "y": 1,
 "z": 1
 },
 "shared_mem": 0,
 "specialize": false,
 "set_launch_bounds": false,
 "max_threads": null,
 "min_blocks_per_sm": 0,
 "specialize_dims": false,
 "passes": "default<03>",
 "codegen_opt": 3,
 "codegen_method": "serial",
 "prune": true,
 "internalize": true
},
```

```
"Result": {
 "preprocess_ir_time": 9.2298723757267e-06,
 "opt_time": 0.006206092890352011,
 "codegen_time": 0.01226967596448958,
 "obj_size": 4792,
 "exec_time": [
 84040,
 82561,
 81761,
 83360,
 76520
],
 "verified": true,
 "executed": true,
 "failed": false,
 "start_time": "",
 "end_time": "",
 "gpu_id": 0,
 "const_mem_usage": -1,
 "local_mem_usage": 0,
 "reg_usage": 12,
 "error": ""
}
```

# Replay a single Kernel

1

## Replay a single kernel invocation

```
> mneme replay \
 -rdb record-example-dir/<static-hash>.json \
 -rid <dynamic-hash> "default<03>"
```



*These parameters can be modified*

By forming valid configuration ranges of these parameters one can search the space and tune the application in respect to some quantity of interest



*Several quantity of interest are supported*

- Execution Time (exec\_time)
- Register Usage (reg\_usage)
- Binary Size (obj\_size)


2

Execution emits a key-value dictionary describing various metrics


```
"Replay-config": {
 "grid": {
 "x": 40000,
 "y": 1,
 "z": 1
 },
 "block": {
 "x": 256,
 "y": 1,
 "z": 1
 },
 "shared_mem": 0,
 "specialize": false,
 "set_launch_bounds": false,
 "max_threads": null,
 "min_blocks_per_sm": 0,
 "specialize_dims": false,
 "passes": "default<03>",
 "codegen_opt": 3,
 "codegen_method": "serial",
 "prune": true,
 "internalize": true
},
```

```
"Result": {
 "preprocess_ir_time": 9.2298723757267e-06,
 "opt_time": 0.006206092890352011,
 "codegen_time": 0.01226967596448958,
 "obj_size": 4792,
 "exec_time": [
 84040,
 82561,
 81761,
 83360,
 76520
],
 "verified": true,
 "executed": true,
 "failed": false,
 "start_time": "",
 "end_time": "",
 "gpu_id": 0,
 "const_mem_usage": -1,
 "local_mem_usage": 0,
 "reg_usage": 12,
 "error": ""
}
```

# How do you go from a single replay to a feedback loop (autotune)?



# How do you go from a single replay to a feedback loop (autotune)?



# Execution With Multiple Workers

1

Create the sampling strategy

2


Invoke submit that returns a future

3

Get results (blocking call)

```
167 SS = ExhaustiveSamplingStrategy(space)
168
169 for i, config in enumerate(SS):
170 if not config.is_valid():
171 continue
172 futures.append((config, [executor.submit(config)]))
173
174 for i, (config, future) in enumerate(futures):
175 val = [future.result()]
```

# Mneme is extremely efficient in performing the feedback loop



## Traditional Benchmark-Centric Measurement

- Full **host + device compilation** per experiment
- Whole-application execution (fork/exec, runtime overheads)
- Repeated **device memory initialization**
- Heavy **I/O and data marshaling**
- Output validation via **separate runs / subprocesses**




## This Breaks Data-Driven Optimization

- Iteration cost measured in **minutes**
- Experiment throughput Repeated **device memory initialization**
- Exploration space collapses prematurely
- Decisions become **sample-starved**

Data-driven approaches don't need hundreds of samples — they need thousands samples on 100s of benchmarks/kernels

# Mneme is extremely efficient in performing the feedback loop



## Traditional Benchmark-Centric Measurement

- Full **host + device compilation** per experiment
- Whole-application execution (fork/exec, runtime overheads)
- Repeated **device memory initialization**
- Heavy **I/O and data marshaling**
- Output validation via **separate runs / subprocesses**

## Mneme Approach

- Compile **only the device code** of the kernel
- Execute only the kernel under investigation
- Single device memory initialization
- Minimal I/O
  - User should use the robust python ecosystem for persistent storage
- Automated bit wise exact validation

Data-driven optimization is fundamentally throughput-limited

# Simple results on MiniFE

## ➤ Time to build miniFE:

- 10s clean build, 4 seconds modifying single source file

## ➤ Time to execute miniFE:

- No Recording : 36 seconds
- With Recording: 38 seconds
  - This cost is paid once
  - Size the GPU snapshot and speed of IO define slowdown

## ➤ Back-of-the-envelope calculation:

- To run 200 experiments and optimize a single kernel, we would need roughly:
  1. Run MiniFE :  $200 * (4 \text{ (compile-time)} + 7 \text{ (number of experiments to reduce noise)} * 36) = 51200 \text{ seconds}$   
**= 0.004 observations/second**
  2. Use sub process + standalone replay tool:  $38 + 200 * (7 \text{ seconds}) = 1438 \text{ seconds} = 0.13 \text{ observations/second}$  ← SC-23
  3. Use python mneme interface (single worker):  $38 + 120 \text{ (seconds)} = 158 \text{ seconds} = 1.26 \text{ observations/second}$  ← Mneme

# Conclusions

➤ **LLVM enables deep GPU optimization — but experimentation cost limits exploration**

- Traditional benchmark-centric workflows are too slow for data-driven optimization
- Experiment throughput — not compiler capability — becomes the bottleneck

➤ **Proteus makes LLVM JIT specialization practical**

- Low-overhead, programmable LLVM JIT for device code
- Specialization and optimization close to the compiler pipeline

➤ **Mneme enables scalable, data-driven GPU optimization**

- Record-replay decouples kernels from full applications
- Orders-of-magnitude faster optimization feedback loops
- Python-driven autotuning and analysis at scale

👉 High-throughput record-replay + LLVM JIT turns GPU optimization into a data-driven workflow